

R Notebook

ECO518 PS4

0. Set up

```
# Load packages
if(!require(pacman)) install.packages("pacman")

## Loading required package: pacman

pacman::p_load(ggplot2, dplyr, sf, tigris,
               viridis, patchwork, sandwich, nlme, jtools, estimatr,
               stargazer, car)
theme_set(theme_bw())
# Set paths
dir <- paste0("/Users/tombearpark/Documents/princeton/1st_year/",
              "term2/ECO518_Metrics2/sims/exercises/4_grouped_data/")
out <- paste0(dir, "out/")
# Load in the data
load(paste0(dir, "caschool.RData"))
df <- tibble(caschool)
# load a shapefile for maps
cal <- counties(state = "California", cb = TRUE)
```

```
## |
```

```
|
```

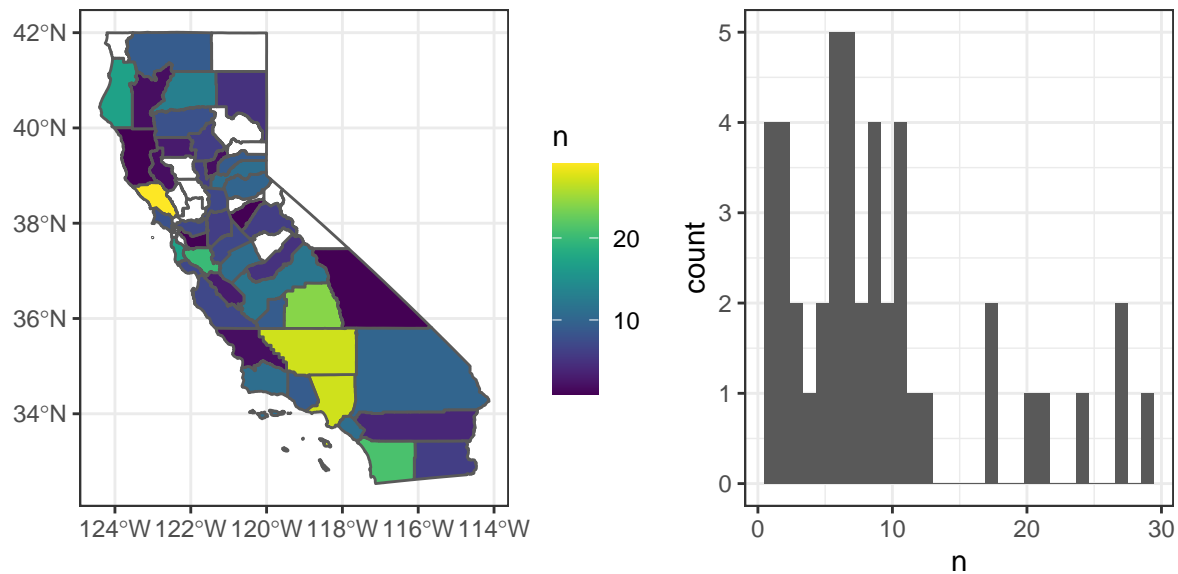
```
plot_df <-
  left_join(cal, df %>% group_by(county) %>% tally(),
            by = c("NAME" = "county")) %>%
  mutate(obs = ifelse(is.na(n), 0, n))

# Make sure the merge worked properly
stopifnot(
  dim(anti_join(df %>% group_by(county) %>% tally(),
                cal, by = c("county" = "NAME")))[1] == 0)

wrap_elements(
  (ggplot(plot_df) + geom_sf(aes(fill = n)) +
   scale_fill_viridis(na.value = "white")) +
  (ggplot(plot_df) + geom_histogram(aes(x = n))))
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

```
## Warning: Removed 13 rows containing non-finite values (stat_bin).
```



2. Exercise

Problem 1

Estimate a linear regression of the average test score (*testscr*) on student-teacher ratio, computers per student, and expenditures per student. Determine whether the three variables have explanatory power by an *F*-Test of the hypothesis that all three have zero coefficients and via the Bayesian information criterion (*BIC*). The latter can be computed from an *F*-statistic: The *BIC* rejects the restriction when the *F*-statistic exceeds the log of the sample size.

```
N <- length(df$avginc)
reg1 <- "testscr ~ str + comp_stu + expn_stu"
lm1 <- lm(data = df, formula(reg1))
```

The *F* stat is 14.96.

```
summary(lm1)
```

```
##
## Call:
## lm(formula = formula(reg1), data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -49.660 -14.093  -0.733  13.079  45.975
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept) 663.317921 19.348195 34.283 < 2e-16 ***
## str         -1.303902  0.606981 -2.148  0.0323 *
## comp_stu    63.638660 14.477669  4.396  1.4e-05 ***
## expn_stu     0.001468  0.001799  0.816  0.4151
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 18.17 on 416 degrees of freedom
## Multiple R-squared:  0.09738,    Adjusted R-squared:  0.09087
## F-statistic: 14.96 on 3 and 416 DF,  p-value: 2.892e-09
```

```
compare_models <- function(lm_restricted, lm_unrestricted, N){
  RSSR <- sum(lm_restricted$residuals^2)
  RSSU <- sum(lm_unrestricted$residuals^2)
  k <- length(lm_unrestricted$coefficients) - length(lm_restricted$coefficients)

  Fstat <- ((RSSR - RSSU) / k) / (RSSU / (N-k-1))
  pVal <- pf(Fstat, k, N-k-1, lower.tail = FALSE)

  BIC_R <- N * log(RSSR / N) + length(lm_restricted$coefficients) * log(N)
  BIC_U <- N * log(RSSU / N) + length(lm_unrestricted$coefficients) * log(N)

  lowerBIC <- ifelse(BIC_R < BIC_U, "restricted", "unrestricted")

  return(
    tibble(Fstat = Fstat, pVal = pVal,
            BIC_R = BIC_R, BIC_U = BIC_U, logN = log(N),
            lowest_BIC_model = lowerBIC))
  }
lm0 <- lm(data = df, testscr ~ 1)
comparison_1 <- compare_models(lm0, lm1, N)
comparison_1
```

```
## # A tibble: 1 x 6
##   Fstat      pVal BIC_R BIC_U logN lowest_BIC_model
##   <dbl>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1  15.0 0.00000000289 2481. 2456.  6.04 unrestricted
```

- BIC is smallest for the true model. BIC is smallest for the more complex model
 - We can also see that the F-stat is larger than the log of the sample size, so we reject the restriction
- F stat strongly rejects the Null that the coefficients aren't jointly significant

```
Anova(lm1)
```

```
## Anova Table (Type II tests)
##
## Response: testscr
##           Sum Sq Df F value    Pr(>F)
## str         1523  1  4.6147  0.03228 *
## comp_stu     6377  1 19.3217 1.404e-05 ***
```

```
## expn_stu      220    1  0.6655   0.41510
## Residuals 137298 416
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# run a clustered version
```

```
lm1_c <- lm_robust(formula(reg1), data = df, cluster = county)
summary(lm1_c)
```

```
##
## Call:
## lm_robust(formula = formula(reg1), data = df, clusters = county)
##
## Standard error type: CR2
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)    CI Lower CI Upper    DF
## (Intercept) 663.317921  21.950633 30.2186 1.273e-21 618.174063 708.46178 25.72
## str         -1.303902   0.656269 -1.9868 5.796e-02 -2.655260   0.04746 25.09
## comp_stu     63.638660  18.873940  3.3718 2.384e-03 24.809323 102.46800 25.55
## expn_stu      0.001468   0.002717  0.5402 5.943e-01 -0.004154   0.00709 22.91
##
## Multiple R-squared:  0.09738 , Adjusted R-squared:  0.09087
## F-statistic: 6.537 on 3 and 44 DF,  p-value: 0.0009409
```

Problem 2

Do the same thing with a regression that adds the demographic variables: Average income, subsidized meals, calWorks per cent, and English learners percent. Again check whether the three "policy variables have explanatory power using an F test and BIC. Here you may need to extract the covariance matrix of coefficients from the lm() output to construct the F or chi-squared statistic.

```
reg2 <- paste0(reg1, " + avginc + meal_pct + calw_pct + el_pct")
lm2 <- lm(data = df, formula(reg2))
```

```
compare_models(lm_restricted = lm1, lm_unrestricted = lm2, N = N)
```

```
## # A tibble: 1 x 6
##   Fstat      pVal BIC_R BIC_U logN lowest_BIC_model
##   <dbl>    <dbl> <dbl> <dbl> <dbl> <chr>
## 1  387. 1.36e-138 2456. 1827.  6.04 unrestricted
```

```
summary(lm2)
```

```
##
## Call:
## lm(formula = formula(reg2), data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -31.0536  -5.3377   0.1755   5.0343  26.4272
```

```
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  6.596e+02  9.023e+00  73.098 < 2e-16 ***
## str          -1.899e-01  2.835e-01  -0.670  0.5034
## comp_stu      1.189e+01  6.898e+00   1.724  0.0855 .
## expn_stu      1.526e-03  8.917e-04   1.712  0.0877 .
## avginc        6.217e-01  8.772e-02   7.087 5.98e-12 ***
## meal_pct     -3.756e-01  3.589e-02 -10.465 < 2e-16 ***
## calw_pct     -7.782e-02  5.722e-02  -1.360  0.1745
## el_pct       -1.981e-01  3.323e-02  -5.962 5.37e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.391 on 412 degrees of freedom
## Multiple R-squared:  0.8093, Adjusted R-squared:  0.806
## F-statistic: 249.7 on 7 and 412 DF,  p-value: < 2.2e-16
```

```
Anova(lm2)
```

```
## Anova Table (Type II tests)
##
## Response: testscr
##           Sum Sq Df F value    Pr(>F)
## str          31.6  1  0.4486  0.50337
## comp_stu     209.2  1  2.9710  0.08552 .
## expn_stu     206.3  1  2.9302  0.08769 .
## avginc      3536.7  1 50.2267 5.979e-12 ***
## meal_pct    7711.7  1 109.5178 < 2.2e-16 ***
## calw_pct     130.3  1  1.8498  0.17455
## el_pct      2502.8  1 35.5437 5.365e-09 ***
## Residuals 29011.1 412
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Once again, our tests prefer the more complex model

```
lm2_c <- lm_robust(formula(reg2), data = df, cluster = county)
summary(lm2_c)
```

```
##
## Call:
## lm_robust(formula = formula(reg2), data = df, clusters = county)
##
## Standard error type: CR2
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|) CI Lower CI Upper DF
## (Intercept) 659.587061 12.566652 52.4871 1.446e-27 633.734800 685.439322 25.57
## str         -0.189910  0.325244 -0.5839 5.645e-01 -0.859845  0.480025 24.94
## comp_stu     11.890284  7.657159  1.5528 1.328e-01 -3.864839 27.645406 25.48
## expn_stu      0.001526  0.001306  1.1685 2.550e-01 -0.001181  0.004234 22.20
```

```
## avginc      0.621673    0.100062    6.2129 3.462e-05    0.405035    0.838311 12.73
## meal_pct   -0.375618    0.044597   -8.4226 4.652e-08   -0.468557   -0.282679 20.30
## calw_pct   -0.077818    0.063414   -1.2271 2.459e-01   -0.217716    0.062080 10.79
## el_pct     -0.198137    0.038746   -5.1137 6.922e-05   -0.279448   -0.116826 18.29
##
## Multiple R-squared:  0.8093 ,    Adjusted R-squared:  0.806
## F-statistic: 407.6 on 7 and 44 DF,  p-value: < 2.2e-16
```

Problem 3

Repeat the previous estimations and tests in models that add county fixed effects. In R using `lm()`, this is accomplished by just adding “county” to the list of right-hand side variables. (county is a “factor” in the R dataframe, so R automatically converts it into the appropriate array of dummy variables when including it in a regression.)

```
reg3 <- paste0(reg2, "+ county")
lm3 <- lm(data = df, formula(reg3))
lm3_c <- lm_robust(data = df, formula(reg3), cluster = county)
```

```
compare_models(lm2, lm3, N)
```

```
## # A tibble: 1 x 6
##   Fstat      pVal BIC_R BIC_U  logN lowest_BIC_model
##   <dbl>      <dbl> <dbl> <dbl> <dbl> <chr>
## 1  2.27 0.0000199 1827. 1994.  6.04 restricted
```

- This time, our model doesn’t want us to choose the extra complexity in the BIC.
- However, our F-stat approach prefers the more complex model

```
Anova(lm3)
```


```
## Anova Table (Type II tests)
##
## Response: testscr
##           Sum Sq Df F value    Pr(>F)
## str          14.7  1  0.2363  0.627173
## comp_stu     133.7  1  2.1483  0.143580
## expn_stu       9.2  1  0.1477  0.700940
## avginc      3149.0  1 50.6039 5.937e-12 ***
## meal_pct    6168.1  1 99.1184 < 2.2e-16 ***
## calw_pct     98.1  1  1.5761  0.210119
## el_pct       596.1  1  9.5798  0.002118 **
## county     6110.7 44  2.2317 3.170e-05 ***
## Residuals 22900.4 368
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
vars <- c("str", "comp_stu", "expn_stu", "avginc", "meal_pct", "calw_pct", "el_pct" )
plot_summs(lm1, lm2, lm3, coefs = vars, scale = TRUE)
```

```
## Registered S3 methods overwritten by 'broom':
##   method      from
##   tidy.glht    jtools
##   tidy.summary.glht jtools

## Loading required namespace: broom.mixed

## Registered S3 method overwritten by 'broom.mixed':
##   method      from
##   tidy.gamlss broom
```



Problem 4


The districts vary greatly in size. Average scores might have more sampling variation in small districts. Plot the squared residuals from the estimated model in 2 against the total enrollment variable. Estimate a linear regression of these squared residuals on $1/\text{enrl_tot}$. Use the inverse of these predicted values as the weights argument in `lm()` (or otherwise estimate the corresponding weighted regression estimates) in the question 2 regression.

First, let's plot the squared residuals against the total enrollment variable.

```
df$u2 <- lm2$residuals
df$sqr_u2 <- lm2$residuals ^ 2
ggplot(df, aes(x = enrl_tot, y = sqr_u2)) +
  geom_point(alpha = .7) +
```

```
xlab("Total Enrollment") + ylab("Squared Residuals") +
geom_smooth(method = lm)
```


```
## 'geom_smooth()' using formula 'y ~ x'
```



Lets also plot the squared residuals against the inverse of total enrollment, since that relationship is what we are going to use for our weighting scheme.

```
ggplot(df, aes(x = 1/enrl_tot, y = sqr_u2)) +
  geom_point(alpha= .7) +
  xlab("Inverse of Total Enrollment") + ylab("Squared Residuals") +
  geom_smooth(method = lm)
```

```
## 'geom_smooth()' using formula 'y ~ x'
```

We can estimate a linear regression, to calculate the weights...

```
df$inv_enrl_tot <- 1 / df$enrl_tot
lm4_w <- lm(data = df, sqr_u2 ~ inv_enrl_tot)
```

Next we run a regression weighted by the inverse of the residuals


```
df$weights_lm4 <- 1 / lm4_w$fitted.values
lm4 <- lm(data = df, formula(reg2), weights = df$weights_lm4)
summary(lm4)
```

```
##
## Call:
## lm(formula = formula(reg2), data = df, weights = df$weights_lm4)
##
## Weighted Residuals:
##      Min       1Q   Median       3Q      Max
## -2.7312 -0.6772  0.0079  0.6082  3.2024
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  6.602e+02  9.052e+00  72.934  < 2e-16 ***
## str         -3.455e-01  2.811e-01  -1.229   0.2198
## comp_stu     1.136e+01  6.986e+00   1.626   0.1048
## expn_stu     2.007e-03  8.895e-04   2.256   0.0246 *
## avginc       6.221e-01  8.106e-02   7.675  1.21e-13 ***
```

```
## meal_pct      -3.857e-01  3.463e-02 -11.138  < 2e-16 ***
## calw_pct      -6.986e-02  5.104e-02  -1.369   0.1718
## el_pct        -1.773e-01  3.078e-02  -5.761  1.64e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.009 on 412 degrees of freedom
## Multiple R-squared:  0.8435, Adjusted R-squared:  0.8409
## F-statistic: 317.3 on 7 and 412 DF,  p-value: < 2.2e-16
```

Compare model 2 to this weighted version...

```
plot_summs(lm2, lm4, scale = TRUE, model.names = c("Unweighted", "Weighted"))
```




Problem 5

For at least two of the above regression models, calculate standard errors clustered by county. This is done very easily with the `vcovCL()` function from the `sandwich` package — so easily that if you're doing it this way you might want to see how much difference it makes in all of the above regressions.

```
plot_summs(lm2, lm2, scale = TRUE,
            robust = list(FALSE, c(cluster = df$county)),
            model.names = c("OLS SEs", "County-Clustered SEs"))
```

```
## Warning in if (type == TRUE) {: the condition has length > 1 and only the first
## element will be used
```



This is done in code throughout. General comment - clustering doesn't make much difference in this setup.

Problem 6

Estimate a random effects model, with county effects. In R, use the `lme()` function from the `nlme` package to estimate the 7-variable regression, with random effects by county. You do this by giving `lme` the argument `random = ~1 | county`. Also use the argument `method="ML"`, so that the estimation is by maximum likelihood.

```
RE <- lme(data = df, formula(reg2), random = ~1 | county, method = "ML")
summary(RE)$tTable
```

	Value	Std.Error	DF	t-value	p-value
## (Intercept)	661.036834029	9.0964677878	368	72.6696174	2.445055e-220
## str	-0.196646628	0.2912970412	368	-0.6750725	5.000536e-01
## comp_stu	12.606870943	6.8619471986	368	1.8372148	6.698437e-02
## expn_stu	0.001068223	0.0008875184	368	1.2036069	2.295152e-01
## avginc	0.664923420	0.0925252296	368	7.1864012	3.727999e-12
## meal_pct	-0.367437398	0.0370049535	368	-9.9294111	9.623245e-21
## calw_pct	-0.084161241	0.0589872135	368	-1.4267709	1.544938e-01
## el_pct	-0.186981628	0.0347232272	368	-5.3849150	1.294861e-07

```
summary(RE)
```

```
## Linear mixed-effects model fit by maximum likelihood
```

```
## Data: df
##      AIC      BIC    logLik
## 2985.113 3025.516 -1482.557
##
## Random effects:
## Formula: ~1 | county
##      (Intercept) Residual
## StdDev:      2.466434 8.004958
##
## Fixed effects: formula(reg2)
##              Value Std.Error DF t-value p-value
## (Intercept) 661.0368  9.096468 368 72.66962  0.0000
## str          -0.1966  0.291297 368 -0.67507  0.5001
## comp_stu      12.6069  6.861947 368  1.83721  0.0670
## expn_stu       0.0011  0.000888 368  1.20361  0.2295
## avginc        0.6649  0.092525 368  7.18640  0.0000
## meal_pct     -0.3674  0.037005 368 -9.92941  0.0000
## calw_pct     -0.0842  0.058987 368 -1.42677  0.1545
## el_pct       -0.1870  0.034723 368 -5.38492  0.0000
## Correlation:
##      (Intr) str      cmp_st expn_s avginc ml_pct clw_pc
## str      -0.910
## comp_stu -0.127  0.142
## expn_stu -0.750  0.513 -0.137
## avginc   -0.097  0.032 -0.024 -0.294
## meal_pct -0.097 -0.010 -0.059 -0.110  0.501
## calw_pct  0.073 -0.006  0.118 -0.129 -0.034 -0.625
## el_pct    0.062 -0.032  0.157  0.005 -0.165 -0.649  0.293
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.86054173 -0.60780373 -0.01266389  0.59529612  2.87465042
##
## Number of Observations: 420
## Number of Groups: 45
```

Problem 7

Compare the random effects 7-variable model to the fixed effects model. In R, you can do this by re-estimating the fixed-effect model with the `gls()` function from the `nlme` package, again being sure to use `method="ML"` argument. The `summary()` function applied to either random effects or fixed effects models computed this way deliver both log likelihood and BIC values, so the models can be compared both by frequentist chi-squared test based on the log likelihood and via the BIC.

```
FE <- gls(data = df, formula(paste0(reg2, "+ county")), method = "ML")
summary(FE)$tTable[1:8, ]
```

```
##              Value Std.Error t-value p-value
## (Intercept) 6.714925e+02 1.288580e+01 52.1110365 5.977392e-172
## str        -1.524921e-01 3.136940e-01 -0.4861174 6.271733e-01
## comp_stu     1.045821e+01 7.135233e+00  1.4657140 1.435801e-01
## expn_stu     3.512003e-04 9.137477e-04  0.3843515 7.009399e-01
## avginc       7.510105e-01 1.055733e-01  7.1136391 5.937188e-12
```

```
## meal_pct      -3.913137e-01 3.930501e-02 -9.9558217  7.808023e-21
## calw_pct      -8.233494e-02 6.558305e-02 -1.2554301  2.101192e-01
## el_pct        -1.226652e-01 3.963174e-02 -3.0951246  2.117931e-03
```

Problem 8

Finally, for the random effects model, use a regression of its squared residuals on $1/(\text{total enrollment})$ to generate weights for a weighted random effects estimation; see if this improves likelihood and/or changes important estimates. [Note: I think that in the nlme estimation functions the “weights” arguments are variance scales — the inverse of the weights used in `lm()`. So you would use a `weights = ~w` argument to `lme()` if you used `weights = 1/w` in `lm()`].

```
df$re_sqr_resid <- RE$residuals[, "county"]^2
lm8_w          <- lm(data = df, re_sqr_resid ~ inv_enrl_tot)
df$re_weights  <- lm8_w$fitted.values

RE_w <- lme(data = df, formula(reg2), random = ~1 | county,
            method = "ML", weights = ~re_weights)
summary(RE_w)$tTable
```

##		Value	Std.Error	DF	t-value	p-value
##	(Intercept)	660.917809859	9.1503020743	368	72.229070	1.979801e-219
##	str	-0.353137974	0.2901083886	368	-1.217262	2.242845e-01
##	comp_stu	12.638742122	6.9264918214	368	1.824696	6.885753e-02
##	expn_stu	0.001556244	0.0008818318	368	1.764785	7.842936e-02
##	avginc	0.686849009	0.0843735423	368	8.140573	6.152343e-15
##	meal_pct	-0.370073550	0.0352340979	368	-10.503279	9.580776e-23
##	calw_pct	-0.075814218	0.0525721136	368	-1.442099	1.501248e-01
##	el_pct	-0.172368003	0.0319298944	368	-5.398327	1.208316e-07


Problem 9

Be ready to discuss: Does the evidence favor an important effect from the “controllable” variables? The sizes and signs of the estimated effects, not just the significance levels of tests, should inform your views on this.


```
comp_df <- data.frame(
  df,
  model_1 = lm1$fitted.values,
  model_2 = lm2$fitted.values,
  model_3 = lm3$fitted.values)

(ggplot(data = comp_df) +
  geom_point(aes(x = model_1, y = model_2, color = avginc)) +
  scale_color_viridis() +
  geom_abline(slope = 1, color = "red") +
  ggtitle("Comparison of fitted values") +
  theme(legend.position = "none")) + (
ggplot(data = comp_df) +
  geom_point(aes(x = model_2, y = model_3, color = avginc)) +
  scale_color_viridis() +
```

```
geom_abline(slope = 1, color = "red") +
ggtitle("Comparison of fitted values"))
```



```
ggplot(data = comp_df) +
  geom_point(aes(x = model_2, y = model_3, color = avginc)) +
  scale_color_viridis() +
  geom_abline(slope = 1, color = "red") +
  ggtitle("Comparison of fitted values")
```



```
library(caret)
```

```
## Loading required package: lattice
```

```
##
```

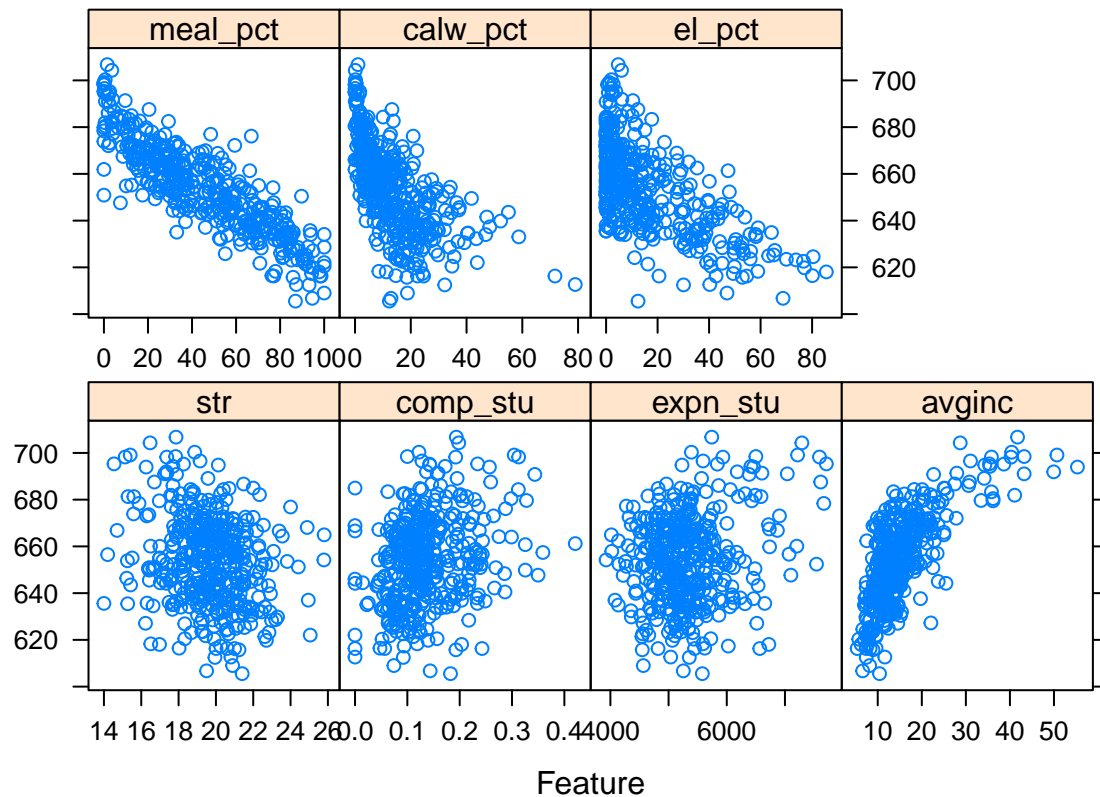
```
## Attaching package: 'caret'
```

```
## The following object is masked _by_ '.GlobalEnv':
```

```
##
```

```
## compare_models
```

```
featurePlot(x = df[vars], y = df$testscr)
```



```
# try a lasso regression
library(glmnet)


## Loading required package: Matrix

## Loaded glmnet 4.1

y <- df$testscr %>% as.matrix()
X <- model.matrix(formula(paste0(reg2, "+ log(avginc)")), df)

# Helped function for processing results
get_results = function(fit){
  tmp_coefs <- coef(fit)
  myResults <- data.frame(
    feature = tmp_coefs@Dimnames[[1]][ which(tmp_coefs != 0 ) ], #intercept included
    coef    = tmp_coefs [ which(tmp_coefs != 0 ) ] #intercept included
  )
  return(myResults)
}

fit <- cv.glmnet(X, y)
plot(fit)
```

```
get_results(fit)
```

```
##      feature      coef
## 1 (Intercept) 665.4959536
## 2      avginc   0.5428310
## 3     meal_pct -0.3909220
## 4      el_pct  -0.1380925
```